



RPS COHORTS • WORKSHOP SUMMARY • 8 AUGUST 2026


From a client brief to a page you would ship

One real engagement, start to finish. Six documents, six prompts, one tool — and the reasoning that connects them. This is the short version: what each stage is for, what it was asked, and what it produced.

THE SHORT VERSION

What this session was

A client brief arrived as an email. Ninety minutes later there was a homepage running in a browser. Nothing in between was improvised — the work moved through six documents, each one the controlling input to the next.




THE ARGUMENT

The quality of what a model builds is set almost entirely by what it was given. So the work is not prompting — it is producing the documents a competent studio would have produced anyway, in an order where each one answers the questions the next one would otherwise invent answers to.

Two properties hold the chain together. **Requirement IDs become a shared vocabulary**, so the research, the audit, the outline and the build prompt all cite the same codes. And **open questions are carried, never resolved silently** — they travel from the brief all the way into the built page, where they appear as visible placeholders rather than plausible inventions.

The five moves

The same craft you already run in Figma, translated move for move. Nothing here is new — only the artefact each move produces.



THE MOVE	IN FIGMA YOU WOULD	IN THIS SESSION
Define	Discovery — brief, competitor audit	Requirements, research and audit, agreed before anything is drawn
Structure	Sitemap, wireframe — blocks and order	Content outline — the wireframe, written in words
Style	Moodboard, colour, type, variables	References gathered, then a foundations page built from them
Specify	Mockup and handoff — redlines to the dev	The build prompt — one executable handoff document
Review	Design QA — crit, then iterate	The impeccable audit — critique, then iterate

CONTENTS

What is in here

A summary of the reasoning, with the prompts printed in full. Each stage opens with what it is for and what it produced, then reproduces the prompt exactly as it was run. The documents themselves are long, so they sit behind links.

–	The client, and the brief that started it	SOURCE MATERIAL
–	The pipeline — six stages, one chain	PROCESS MAP
01	Requirements	FULL PROMPT + DOCUMENT
02	Competitive research	FULL PROMPT + DOCUMENT
03	UX audit	FULL PROMPT + DOCUMENT
04	Content outline	FULL PROMPT + DOCUMENT
05	Design system foundations	STYLE, PROMPT, RESULT
06	The build prompt	FULL PROMPT
07	The page it produces	THREE LIVE BUILDS
08	Going live, and what to expect	FIELD NOTES

HOW TO USE THIS

Replace the Tessera-specific substance — the figures, the four pillars, the requirement IDs — with your client's, and keep the scaffolding. The load-bearing rules transfer to any engagement: five passes on the brief, MoSCoW priorities, requirement IDs as a shared vocabulary, and open questions that are carried rather than quietly answered.

SOURCE MATERIAL

The client, and the brief that started it

Tessera — serverless GPU infrastructure for AI applications. Point it at a Dockerfile; it provisions, scales and fails over across clouds. Pay per second.

THE NUMBERS THE PAGE HAS TO CARRY

2-4s	cold starts via memory and GPU snapshotting, against a field sitting at 30+ seconds
12+	GPU varieties, no reservations
99.999%	uptime, multi-region failover
SOC 2	plus HIPAA, GDPR, ISO · gVisor isolation

WHAT THE BRIEF ASKED FOR

Homepage only. Four pillars. Security as a proper section, not a badge strip. Three or four blog posts. 3D graphics — but no spinning GPU chip, and nothing that tanks Lighthouse. Lead capture is non-negotiable, on both the enterprise and self-serve paths, with capture points through the scroll and work email only. Copy is in scope: no lorem ipsum. Two revision rounds. Three weeks to initial directions.

THE READER, WHICH IS REALLY TWO READERS

The audience is technical and commercially sceptical, and it splits: the ML engineer or infra lead who deploys on a self-serve path, and the CTO who signs on an enterprise one. Both are described as allergic to marketing language. They will read the site's own Lighthouse score as a proxy for the platform's credibility.

This is what makes it a B2B page and not a B2C one. The audience is bounded, so one page can serve both without a toggle — but proof has to beat persuasion on every section: numbers, methodology, named customers, certifications.

A NOTE ON REAL BRIEFS

This one was unusually complete. Most arrive as “we like Apple's website”, “monk-like composure”, or a single actor's name. The stage that follows exists precisely to turn that kind of sentence into something two teams can sign.

PROCESS MAP

Six stages, one chain

Every arrow is an attachment in the next prompt. Nothing is retyped, and nothing is summarised on the way through — the previous document goes in whole.




FIG 1 • THE DOCUMENT PIPELINE • EACH STAGE IS ONE PROMPT AND ONE ARTEFACT

Stages 02 and 03 run in parallel and both answer to stage 01. Stage 05 — the design system — was the missing step in the original walkthrough: it is what turns a folder of references into a set of decisions the build prompt can inherit without improvising. Stage 06 consumes all five.

01 Requirements

STAGE 01 · MOVE: DEFINE

The brief is a conversation. This turns it into something two teams can sign. Nothing in the email is allowed to evaporate between the inbox and the build.

DOCUMENT CREATED

02_Requirements.docx – v1.0

[Open in Google Drive ↗](#)

WHAT IT FEEDS

Every stage after this one. The IDs become the shared vocabulary.

01-Prompt-RequirementsGen.md

REPRODUCED IN FULL

You are acting as the lead consultant at a design and technology studio. A client brief has arrived. Convert it into a formal requirements document that can be sent back to the client for sign-off and forward to an engineering team for build.

INPUTS

CLIENT BRIEF:

[Paste the full email, or attach the file. Include the sender's name, role and any cc'd contacts – these become named owners in the document.]

STUDIO: [Your studio name]

STUDIO LEAD: [Name, role – the person who signs]

DATE: [Today's date]

DOCUMENT VERSION: [v1.0 unless revising]

OUTPUT FORMAT: [Word document / Markdown / inline]

YOUR TASK

Do not summarise the brief. Convert it. Every instruction in the brief becomes a numbered, testable requirement. Nothing is dropped, softened, or merged away.

Work through five passes before you write anything.

Pass 1 – Extract the explicit

Go line by line through the brief. Every sentence that constrains, requests, permits or forbids something becomes a requirement. Watch for instructions buried in casual phrasing – "oh and", "obviously", "we'd like", "just", "don't bother with" all carry scope. Preserve the client's exact figures, product names, technical terms and claims verbatim; you may rewrite the sentence around them, never the substance.

Pass 2 – Surface the implied

The brief will state an outcome without stating what the outcome depends on. Name those dependencies as requirements so they are agreed rather than assumed. Typical gaps: states and edge cases for anything interactive, validation and error handling, empty and minimum-content states, fallbacks when a technology is unsupported, accessibility, tracking and measurement, content length constraints, and the performance budget implied by any ambitious visual requirement.

Mark these clearly as studio-added rather than client-stated where it matters.

Pass 3 – Capture the negatives

Anything the client ruled out is as binding as anything they asked for, and is the

thing most often silently reintroduced later. Record every exclusion as its own requirement at Won't priority. This includes both scope exclusions ("not the docs site") and creative exclusions ("not a spinning logo").

Pass 4 – Find the conflicts

Requirements that pull against each other are the real risk in the engagement. Find them, state them plainly, and specify how the trade-off gets resolved and when. Never resolve a client-level conflict silently in the document.

Pass 5 – Find the holes

Identify every decision that cannot be made without client input. Each becomes an open question with the requirements it blocks listed against it. If the brief asks "what else do you need from us", this section is the answer. Include the gap the client did not notice they had left – the missing input that a professional reading the brief would spot immediately.

DOCUMENT STRUCTURE

Adapt to the brief; do not pad sections the material does not support.

1. Purpose of this document – what it is, what MoSCoW priorities mean, ID prefix key
2. Background and context – the client's business, why the work is happening, audience
3. Scope – in scope / explicitly out of scope / noted for later, all as requirements
4. Structure or architecture of the thing being built
5. ... one section per major content or functional area named in the brief ...
- N-8. Non-functional requirements – performance, responsive, accessibility, technical
- N-7. Deliverables
- N-6. Timeline and process
- N-5. Assumptions and client dependencies
- N-4. Open questions
- N-3. Risks
- N-2. Change control
- N-1. Sign-off block

REQUIREMENT FORMAT

Present requirements in tables: ID | Requirement | Priority.

- IDs use a three-letter prefix by area (SCP scope, LED lead capture, VIS visual, PRF performance, etc.) plus a two-digit number. Define the prefix key up front.
- IDs are stable across versions. Superseded requirements are marked superseded, never renumbered or deleted.
- Priority is Must / Should / Could / Won't.
- One requirement per row. If a row contains "and", check whether it is two.
- Each requirement must be testable – a reader must be able to say whether it was met.
- Cross-reference related requirements by ID rather than repeating them.

Client-supplied substance (feature descriptions, compliance claims, specifications) goes in its own reference table, presented as authoritative source, before the requirements derived from it.

WRITING STANDARD

- Neutral, precise, unhedged. This is a contractual document.
- Never invent a fact, figure, date, name or capability not present in the brief.
- Where you propose something the client did not state, say so and mark it for confirmation. Proposals and instructions must be distinguishable at a glance.
- Where the client's own words carry the requirement, quote the substance exactly.
- No filler. A short section that says something beats a long one that does not.
- Add short interpretation notes where a requirement needs its reasoning attached, but keep them clearly separate from the requirements themselves.

AFTER THE DOCUMENT

In your chat reply, briefly cover:

- The judgement calls you made and why
- The single most important open question
- The gap in the brief the client will most want flagged

02

STAGE 02 · MOVE: DEFINE

Competitive research

Two problems drove the engagement, and both are about the market rather than the client: everyone in the category says the same thing, and the visual language of the space is two years stale.

DOCUMENT CREATED

02_Tessera-Competitive-Positioning-Research.docx

[Open in Google Drive ↗](#)

WHAT IT FEEDS

Positioning and copy decisions in the content outline; the reference set for style.

02-Prompt-CompetitiveResearch.md

REPRODUCED IN FULL

You are acting as the research lead on a homepage rebuild for Tessera, a serverless GPU infrastructure platform for AI applications. The attached requirements document (v1.0) is the controlling scope reference – read it fully before starting and treat its requirement IDs as the vocabulary for your findings.

Produce a competitive research report that ends in a set of concrete, defensible recommendations for the Tessera homepage revamp.

Context you must hold throughout

Two problems drove this engagement, both stated by the client:

1. Category sameness. Everyone in this space writes "scale effortlessly" and "built for AI". Tessera reads as one of many.
2. Dated visual language. Flat gradient-and-grid reads as 2023. Category leaders have moved to dimensional, rendered work.

The audience is technical and commercially sceptical, and splits two ways: ML engineers / infra leads on a self-serve path, and CTOs at AI-native companies on an enterprise path. They are described as allergic to marketing language. They will check the site's own Lighthouse score as a proxy for the platform's credibility. Roughly half the sales pipeline is in regulated industries.

Every conclusion you draw must be useful to a designer, a copywriter or an engineer on Monday morning. Interesting-but-unactionable observations do not belong in this report.

Competitor set

Cover these three tiers. Name any competitor you add and say why.

Tier 1 – direct (serverless GPU / inference platforms). Modal, RunPod, Baseten, Replicate, Beam, Cerebrum, Fal, Together AI, Fireworks AI, Koyeb.

Tier 2 – adjacent (GPU cloud and general serverless competing for the same budget). CoreWeave, Lambda, Northflank, Fly.io, Google Cloud Run GPU, AWS SageMaker Serverless Inference.

Tier 3 – craft benchmark (not competitors; the standard the client is measuring against). Vercel, Linear, Resend, Clerk, Stripe, Cloudflare, and any developer-infrastructure site currently doing dimensional 3D well.

For Tier 3, do not analyse positioning. Analyse execution only: what specifically makes the page read as current, and what it costs in page weight and performance to achieve it.

What to analyse – per competitor

A. Positioning and message

- The single-sentence claim in the hero, quoted verbatim.
- The category noun they use for themselves (platform, cloud, runtime, engine).
- Which of the four Tesseract pillars they also claim, and how strongly: cold start, elastic scaling, bring-your-own-code, observability.
- Whether they lead with a number or an adjective.

B. Specificity index

Score each competitor 0-5 on how much verifiable, falsifiable detail appears above the fold and in the first three sections. Count concrete figures, named mechanisms (snapshotting, gVisor, OpenTelemetry, multi-region failover) and real artefacts (code blocks, CLI output, Dockerfiles). This is the single most important comparison in the report, because CPY-02 to CPY-05 turn on it. Show your scoring in a table.

C. Cold start as a competitive claim

Cold-start time is Tesseract's strongest stated differentiator (PIL-02). Build a dedicated comparison: who publishes a cold-start figure, what number, measured how, and where on the page it sits. Flag any competitor claiming faster than 2-4 seconds, and what qualifications sit around that claim. This is the section that most directly determines whether the differentiator survives contact with the market.

D. Page structure and scroll architecture

Section-by-section order for each Tier 1 competitor. Where does trust content sit? Where does the first conversion point sit? How deep is the page? Produce a comparative structural map – this feeds STR-02.

E. Conversion architecture

- Dual-path handling: how do they serve enterprise and self-serve on one page without forcing a wrong choice (LED-04)?
- Number and placement of capture points through the scroll (LED-05).
- Form field count and which fields are mandatory (LED-07, LED-08).
- CTA label language, quoted verbatim – collect these into one list so the pattern is visible (LED-17).
- Where the enterprise path visibly diverges from the developer path.

F. Trust, security and compliance

How is compliance treated – a badge strip, or a real section with explanatory substance? Note exact certification wording used across the set (certified vs compliant vs aligned vs in progress); this informs TRU-07 and CPY-06. Note which competitors surface data residency and workload isolation, and how deep they go.

G. Social proof

Logos, named case studies, usage metrics, testimonials, benchmark posts. Record what each competitor has and where it sits on the page. Tesseract supplied none of this (OQ-07) – this section must quantify the gap so it can be raised with the client with evidence attached.

H. Blog and technical writing as acquisition

Does the homepage surface engineering writing, and with how much design weight? What content types appear (benchmarks, architecture deep-dives, customer builds)? Where does it sit in the scroll (BLG-03, and the placement argument in Section 8)?

I. Visual language

- Flat vs dimensional. Where 3D or rendered work is present: what is it depicting, is it abstract or literal, is it real-time or pre-rendered, is it animated, and does it degrade.
- Where it sits (hero only, or throughout).
- Colour, surface treatment, type, motion language.
- Identify the specific visual moves that have become category clichés and must be avoided (VIS-05, VIS-06). The client has already ruled out a spinning GPU chip; find the rest of the list.

J. Performance benchmark

For every Tier 1 and Tier 3 site, record Lighthouse scores (mobile and desktop) and Core Web Vitals where obtainable, plus approximate page weight and whether heavy 3D assets are lazy-loaded. State your measurement method and its limits. Cross-reference against visual ambition – the point is to find the sites that achieve dimensional craft and pass, and describe how (PRF-01 to PRF-04, VIS-11, VIS-13).

Synthesis

1. Positioning map. Plot the set on two axes: specificity of claim (vague ↔ verifiable) and visual currency (flat ↔ dimensional). Locate Tessera's current position and the target position.
2. The sameness inventory. A list of the phrases, layouts, section patterns and visual devices that recur across three or more competitors. This is the do-not-use list, and it is a direct deliverable against CPY-03.
3. Whitespace. Claims, structures, proof formats or visual territories that nobody in the set currently owns and that Tessera's supplied material can legitimately support.
4. Threat read. Where a competitor's claim directly undercuts one of Tessera's four pillars, say so plainly. Do not soften it.

Recommendations for the homepage revamp

This is the section the report exists for. Make it the longest and the most specific.

1. Positioning recommendation – the claim Tessera should own, with the reasoning from the evidence above, plus two alternates. Written in the register of Section 11.1, not as marketing copy.
2. Structural recommendation – proposed section order with rationale, respecting STR-01 (four pillars, security, blog as distinct sections) and TRU-06 (security reached before the final conversion point). Note where you diverge from the proposed order in Section 4 of the requirements and why.
3. Conversion recommendation – where capture points should sit through the scroll and what each is for, as the first input to the scroll-depth conversion map (LED-06, DEL-06).
4. Copy recommendation – the language territory to occupy and the specific phrases to avoid, drawn from the sameness inventory.
5. Visual recommendation – the 3D territory to pursue and, importantly, the ones to avoid, with the performance evidence from section J attached to each (VIS-11, VIS-13).
6. Proof recommendation – what social proof Tessera needs to compete, ranked by effort against impact, framed as the case to put to the client under OQ-07.
7. Answer to OQ-11 – name the sites the client should look at as references and the ones Tessera should not resemble. This open question is explicitly assigned to research.

Rules

- Quote competitor copy verbatim and briefly where the exact wording is the point; otherwise paraphrase.
- Every claim about a competitor traces to a URL. Note the date you looked.
- Where a figure could not be verified, say so rather than estimating.
- Do not recommend anything that contradicts a Must-priority requirement. If your research suggests one should change, flag it separately as a challenge to the requirement rather than quietly designing around it.
- Tag each recommendation with the requirement IDs it serves.

Output

A Word document, structured with a two-page executive summary at the front that a founder can read on a phone, then the full analysis, then the recommendations. Comparison tables wherever the data supports them. Screenshots where a visual point cannot be made in text.

03

STAGE 03 · MOVE: DEFINE

UX audit

The pre-work, not a general critique. Every finding either justifies a decision the requirements already made, or exposes something they missed.

DOCUMENT CREATED

02_Tessera-homepage-UX-audit.docx

[Open in Google Drive ↗](#)

WHAT IT FEEDS

New requirements back upstream; section-level fixes into the outline.

[Evidence: 00-TesseraOldWebsite.pdf ↗](#)

02-Prompt-UXAudit.md

REPRODUCED IN FULL

You are conducting a UX audit of Tessera's current website, ahead of the homepage rebuild scoped in the attached requirements document (v1.0). The evidence is the attached screenshot PDF. Read the requirements document first – the audit's job is to establish what the new page must fix, in the language the project already uses.

This is the pre-work, not a general critique. Every finding should either justify a decision already made in the requirements, or surface something the requirements missed.

Method

Work through the PDF page by page. Before you start finding problems, build a screen inventory: number every screen or section in the PDF, name it, and describe what it is trying to do. Note where the capture is partial, where scroll position is ambiguous, and where two pages appear to show the same section. The audit references these numbers throughout.

Then audit against this framework. Static screenshots are the only evidence, so state clearly what you can assess and what you cannot.

1. First-encounter and positioning clarity

Can a visitor tell what Tessera is, who it is for and why it is different within one screen? Is the hero claim specific or category-generic? Apply the client's own test: does this read as written by someone who has deployed software, or by someone who has not? Quote the current copy and mark it against the tone reference in Section 11.1.

2. Information architecture and scroll structure

Map the current section order. Assess against STR-01 to STR-06: are the four pillars present with equal weight (PIL-01)? Is security a substantive body section or a footer badge strip (TRU-01, TRU-02, TRU-04)? Does the blog module read as an afterthought (BLG-03)? Where does an enterprise evaluator hit trust content relative to the final CTA (TRU-06)?

3. Content specificity and credibility

Inventory every concrete figure and named mechanism currently on the page – and every place a specific claim has been softened into an adjective. The 2–4 second cold start, 12+ GPU varieties and 99.999% uptime figures are the client's strongest assets (CPY-04); record where they appear, how prominently, and whether the mechanism behind each survives (CPY-05, PIL-05). Flag category-generic phrasing explicitly (CPY-03).

4. Conversion architecture

- Locate every CTA and every form. Record label text verbatim, visual weight, and scroll position.
- Assess the dual-path requirement: are enterprise and self-serve both served, or is one buried (LED-02, LED-03, LED-04)?
- Are capture points distributed through the scroll, or is there one form at the bottom (LED-05)?

- Count form fields and identify which are mandatory (LED-07).

- Identify which form states are visible in the captures and which cannot be assessed – error, loading, success and failure states are frequently absent from a screenshot set, and their absence is itself a finding against LED-11 to LED-13.

5. Trust and risk reduction

What does an enterprise buyer evaluating vendor risk actually get? Assess compliance treatment, data residency, isolation and uptime against TRU-03 and TRU-05. Note the absence of social proof of any kind and its consequence for an enterprise conversion page (OQ-07).

6. Visual design and category currency

Assess against the client's own diagnosis. Is this the flat gradient-and-grid language they describe as reading 2023? Evaluate hierarchy, type scale, spacing discipline, surface treatment, colour use and image quality.

Identify where the visual language actively undermines the technical credibility the copy is trying to establish. Be direct – the client has already stated this position, so hedged language here is unhelpful.

7. Mobile and responsive

Where mobile captures exist, assess whether mobile is a considered design or a reduction of desktop (RSP-02). Pay attention to conversion points, the security block and form ergonomics (RSP-06, LED-19). Where mobile captures are absent, list what needs capturing rather than inferring.

8. Accessibility (as far as static capture allows)

Assessable from screenshots: text contrast against backgrounds including any over imagery or gradient (ACC-02), text size and legibility, touch target size (ACC-07), reliance on colour alone to convey meaning, and link affordance. Not assessable: keyboard operability, focus states, semantic structure, form label association, motion preference handling. List these as an audit gap requiring live-site testing, not as passing findings.

9. Blog module

Assess current treatment against BLG-01 to BLG-08: post count, metadata present, title truncation, thumbnail handling, route to the index, and design weight relative to the rest of the page.

10. Performance signals

Screenshots cannot measure performance, but they can flag risk: heavy imagery, uncompressed assets, large hero media, font loading, layout-shift-prone patterns. Record these as hypotheses to verify against the live site, and note that PRF-01 makes this a hard gate for the rebuild.

Findings format

Every finding gets:

Field	Content
ID	UAX-##
Screen	Reference to the numbered screen inventory
Category	One of the ten areas above
Severity	Critical / High / Medium / Low
Observation	What is on the screen. Description only, no judgement.
Why it matters	The consequence for one of the two named audiences.
Requirement link	The requirement ID(s) in the new scope that address it, or NEW if nothing does.
Recommendation	What the rebuild should do instead.

Severity definitions

- Critical – actively costs qualified leads or destroys credibility with the technical audience.
- High – materially weakens the commercial argument on the page.
- Medium – noticeable friction or inconsistency; fix within normal course of the rebuild.
- Low – polish.

Findings marked NEW are the most valuable output of this audit, because they represent gaps in the requirements document itself. Collect them into their own section at the end and state, for each, whether it should become a new requirement or a change request under Section 21.

Rules

- Separate observation from judgement in every finding. If you cannot point at the pixel, it is not an

observation.

- Do not manufacture findings to fill a category. If section 9 yields two findings, it yields two.
- Do not report anything the screenshots cannot support. Maintain a running list of things that require live-site verification and present it as its own section.
- Where a finding confirms a decision already recorded in the requirements, say so – that is evidence the scope is right, and it is worth having on the record for the client.
- Be specific about copy. Quote it.

Output

1. A findings table covering everything, sorted by severity.
2. A short narrative report: the top five problems, what they cost commercially, and what the rebuild must do differently.
3. The NEW findings section – gaps in the current requirements.
4. The live-site verification list.

Deliver as a Word document with the screenshots embedded next to their findings.

04

STAGE 04 · MOVE: STRUCTURE

Content outline

The wireframe, written in words. It sits between the requirements and the copy deck, and it is the document the client argues with — which is exactly the point.

DOCUMENT CREATED

03_Content_Outline.docx — v0.1

[Open in Google Drive ↗](#)

WHAT IT FEEDS

The build prompt, and the copy deck after it.

03_Content_Outline_Prompt.md

REPRODUCED IN FULL

Tessera — Content Outline Prompt

Reverse-engineered from `Tessera\_Homepage\_Content\_Outline.docx` v0.1. Written to slot in as **Prompt 3** alongside the competitive research and UX audit prompts.

PROMPT 3 — Homepage Content Outline

> **Attach:** `02\_Requirements.docx`
> **Attach if available:** competitive research report (Prompt 1), UX audit (Prompt 2), brand kit
> **Output:** Word document

You are the content architect on a homepage rebuild for **Tessera**, a serverless GPU infrastructure platform for AI applications. Produce the **content outline** for the new homepage.

Read the attached requirements document (v1.0) in full before you write anything. It is the controlling scope reference and its requirement IDs are the vocabulary for this deliverable.

What this document is, and what it is not

This artefact sits between the requirements document and the copy deck. It defines, section by section: what each section is for, what it must contain, how much of it there is, and what commercial work it does.

It is **not** the copy deck (DEL-03), **not** the design, and **not** the conversion map (DEL-06) — though it is the direct input to both. Do not write final copy. Do not describe layouts in visual detail. Stay at the level of content architecture.

Context you must hold throughout

The audience splits two ways and the page must serve both without forcing a wrong choice: ML engineers and infra leads on a self-serve path, and CTOs at AI-native companies on an enterprise path. They are described as allergic to marketing language and will read code snippets as code, not decoration. Roughly half the sales pipeline sits in regulated industries, which is why the security section does real commercial work rather than decorative reassurance.

Two problems drove the engagement: the site reads as category-generic, and its visual language reads as dated.

Your job is the first of those.

Structure of the deliverable

Open with:

1. ****A "how to read this document" preamble.**** State what the artefact is, how indicative copy should be treated, what the character counts are for, and how blocked content is marked.
2. ****A structural summary table**** – every section numbered, with its status (client-specified vs proposed vs conditional), its conversion role, and its requirement anchors. This table is how a reader who opens the document for ninety seconds understands the page.
3. ****The scroll-depth conversion logic**** – where the capture points sit and why, in a short paragraph. This seeds LED-06 and DEL-06 without pre-empting them.

Then the section-by-section body. Then close with a dependency table and a short set of recommendations for kickoff.

Per-section format

Every section gets, in this order:

- ****Purpose**** – one or two sentences on what it is for commercially, not descriptively.
- ****Contents**** – the elements present.
- ****Indicative copy**** – in a monospace block, labelled by slot (`H1`, `Subhead`, `Body`, `CTA`, `Micro`).
- ****Copy constraints**** – a table of character counts and line limits per slot.
- ****Notes**** – the decisions, the risks, and the requirement obligations that bear on the section.
- ****Flags**** – anything blocked, marked against its open question.

Rules for indicative copy

- ****Write real copy, not placeholder.**** CPY-01 rules out lorem ipsum in any delivered artefact, including this one. Every slot gets plausible content in the correct register.
- ****Mark it clearly as indicative.**** It is pre-engineer-call and pre-copy-deck. Say so in the preamble.
- ****Match the register in Section 11.1 of the requirements.**** Declarative, specific, mechanism-forward. Write as someone who has deployed software, for people who have deployed software.
- ****Lead with the client's figures**** – 2-4 second cold starts, 12+ GPU varieties, 99.999% uptime (CPY-04).
- ****Preserve the mechanisms.**** Snapshotting, gVisor, OpenTelemetry, multi-region failover. A figure without its mechanism does not satisfy CPY-05 or TRU-05. State explicitly under each pillar which mechanism must survive to final copy.
- ****Do not use**** "scale effortlessly", "built for AI", "seamlessly", "lightning-fast", "effortless", "enterprise-grade" (CPY-03). Do not use "Request a demo", "Get started today" or "Unlock your free trial" for CTAs (LED-17).
- ****Do not alter client-supplied substance or figures.**** You are writing the surface, not the claim.

Rules for character constraints

Give a range and a line limit for every text slot, at desktop and mobile where they differ. Set them so real content cannot break the layout (CPY-08). State in the preamble that these assume the design system is not yet fixed and tighten at DEL-02.

Rules for structure

- Follow the section order proposed in Section 4 of the requirements as your default.
- The three client-specified sections – four pillars, security, blog – must each be distinct and substantive (STR-01). Security in particular must be a body section, not a footer badge strip (TRU-02, STR-04).
- Security must be reached before the final conversion point (TRU-06).
- Capture points must be distributed through the scroll rather than concentrated at the bottom (LED-05, STR-03). For each, state which audience it is weighted toward and why it sits where it does.
- ****You may propose a change to the order or grouping****, because STR-02 leaves it open. If you do, argue it explicitly in the structural summary, cite the requirement that supports it, and mark it as requiring client agreement. Do not quietly restructure.
- Specify empty and minimum states wherever content is variable – particularly the blog module (BLG-08, STR-06). State the behaviour at each count, including zero.

Rules for blocked content

The requirements carry open questions. Several of them block content directly.

- Mark every blocked item at the point of use, against its open-question ID. Do not only collect them at the back.
- **Do not invent answers.** If certification wording is unconfirmed, write the copy and flag that it cannot go to final until OQ-06 is answered. If no social proof was supplied, do not write fictional logos, metrics or testimonials – say the section is conditional and state what happens if the content does not exist.
- **Distinguish blockers that change structure from blockers that change wording.** Some open questions determine whether a section exists at all, or where it sits. Call these out separately – they need answering at kickoff, not at copy stage.
- Close with a dependency table: reference, blocked content, owner, and what it is needed by.

Where a section's content depends on a visual decision

Note it, briefly, and stop. Where a 3D or motion moment carries meaning, say what conceptual territory is worth exploring and what has become a category cliché, and state the fallback obligations (static fallback, reduced-motion alternative, meaning surviving the asset's absence). Do not art-direct.

Where positioning is not yet settled

The hero claim is a direction-stage decision (DEL-01). Give a primary direction plus two alternates, each with a one-line note on what it wins and what it costs. Do not present a single option as settled.

If the competitive research and UX audit are attached

Use them as evidence, not as content to summarise.

- The **sameness inventory** from the research constrains your copy – nothing you write should appear on it.
- The **whitespace findings** inform the hero direction and the alternates.
- **Critical and High findings** from the audit should each map to a section of this outline. Where one does not, say so.
- Where the audit surfaced **NEW** findings – gaps in the requirements – reflect them here and flag them as needing a change request under Section 21.

Rules

- Tag every section and most notes with the requirement IDs they serve.
- Every recommendation must be actionable by a copywriter, designer or engineer without a follow-up conversation.
- Where you deviate from the requirements, argue it. Where the requirements are silent, say they are silent rather than filling the gap invisibly.
- Do not pad. If a section needs four notes, it gets four.
- Close with a short set of recommendations for kickoff – the things that need a client answer or a decision before the next artefact can start.

Output

A Word document. Cover block with client, source document, version and status. Structural summary table near the front. Indicative copy in visually distinct monospace blocks that do not split across page breaks. Constraint tables per section. Dependency table and kickoff recommendations at the end. Page footer with document name, version and page number.

Notes on using this prompt

It assumes the requirements document exists and is detailed. Against a thinner brief it will over-produce – flagging blockers that were never going to be answered and specifying constraints with no basis. Loosen the character-count and empty-state instructions in that case.

The "do not invent answers" rule is the load-bearing one. Without it, a model asked for a content outline will cheerfully populate a social proof section with plausible-sounding customer logos, and the omission becomes invisible to the client rather than raised with them.

****The structure/wording blocker distinction is worth keeping**** for any engagement with open questions. It is what turns a list of unknowns into a kickoff agenda.

05

STAGE 05 • MOVE: STYLE • PART ONE

Deciding the style

Before a single token is chosen, the direction is gathered. This is the moodboard move — the same one you already run — except the output is a list you can hand to a model rather than a board only you can read.

WHERE	SOURCE	WHAT IT IS ACTUALLY FOR
MOTION & INTERACTION	motion.dev	The animation library and, more usefully, its documentation of what good motion values look like.
	emilkowal.ski	Emil Kowalski on interface animation. Read for judgement, not for code.
AWARD GALLERIES	Awwwards	For the ceiling of the craft. Filter hard — most of it is unshippable.
	Godly	Better curated than most. Good for dark, dimensional, product-led work.
COMPONENT SOURCES	21st.dev	Component patterns you can lift the structure of.
	React Bits	Small interaction pieces. Useful for one moment, not a whole page.
	shadcn/ui	The baseline the audience already recognises. Deviate deliberately.
PRODUCT & PATTERN LIBRARIES	Mobbin	Real product screens. The paid tool that organises; use it to answer specific questions.
	Land-book	Landing pages by category. Good for structure, weak on interaction.
	Lapa	Similar. Cross-check against Land-book rather than trusting either.
	Pinterest	Free, disorganised, occasionally the best thing you find all week.

THE FOUR-LINE REFERENCE TEST

A reference earns its place only if you can write four lines about why it is here: what specifically you are taking, what you are leaving, which section it applies to, and what it would cost you if you copied it whole. A screenshot with no sentence attached is decoration.

PAID VS FREE

Paid tools organise — Mobbin answers a specific question quickly. Free tools inspire — Pinterest and the galleries are where the unexpected thing turns up. Use the first when you know what you are looking for, the second when you do not.

STAGE 05 · MOVE: STYLE · PART TWO

The prompt

One prompt, one file out. The brand colours and typefaces go in fixed; the references set the direction; the eight sections are named so nothing is left to taste at the wrong moment.

05-design-system-foundations.txt

REPRODUCED IN FULL

PROMPT 05 — DESIGN SYSTEM FOUNDATIONS

Attach: 02_Requirements.docx (v1.0), 03_Content_Outline.docx (v0.1), the competitive research report, and the reference screenshots collected above.

Output: one self-contained HTML file — 04-Design System Foundations.html

You are the design systems lead on the Tessera homepage rebuild. Produce the visual foundations: the smallest set of decisions every later screen inherits. Not a component library, not the page itself. The layer underneath both.

THE BRAND INPUT — FIXED, DO NOT REDESIGN

Palette, exactly four colours. Nothing is added:

Void	#0A061A	base; reads as black, but there is purple in it
Brand	#A855F7	the primary purple
Magenta	#B440CB	deeper accent
Paper	#FFFFFF	text

Typefaces, three:

Outfit	headings and display
Schibsted Grotesk	body copy, tables, anything read at length
JetBrains Mono	labels, specs, code

Everything else — every surface, stroke, muted text, every state — is derived from those four colours as an alpha of Paper or a tint of Brand. If you find yourself reaching for a fifth hex value, you have made an error upstream.

THE DIRECTION — FROM THE REFERENCES

Dark-first: purple in the black, not purple on black. Depth comes from one ambient halo and translucent surfaces over it, never from decorative gradient washes. The category reads flat and 2023; this reads dimensional and lit. Restraint is the differentiator — the audience is technical and will read excess as a lack of confidence.

WHAT YOU MUST DECIDE AND STATE

Every value is stated as a number and shown as the thing it describes. A page that describes a shadow without casting one has failed. For each of the eight sections below, render the artefact, then print its specification in mono directly beside it.

01 PALETTE — four swatch cards, 150px tall, each with name, hex, and one line on what it is for. Beneath them, a row of derived token chips: text/primary 100% white, text/secondary 62%, text/muted 42%, stroke/hairline white 10%, surface/raised white 6%.

02 TYPE — a specimen ladder: Display (Outfit 600, 64/66, -3.5% tracking), Heading one (500, 44/50), Heading two (500, 30/36), Heading three (500, 21/28), body (Schibsted Grotesk 17/1.62), small (15/24), overline (mono 12, 0.16em, uppercase). Set a rules panel beside it: Outfit only above 20px, below that its geometry muddies; negative tracking scales with size; weights are 300 for oversized display, 500 default, 600 for emphasis,

never 700 in UI; table numerals use tabular figures. Close with a figure specimen.

03 GRADIENTS – exactly six, named, with the CSS declaration in mono and one line on where each is allowed: Brand (135deg, brand to magenta – primary buttons, active states); Brand extended (115deg, four stops from lilac through to deep plum – hero fills that need internal light); Depth (180deg, #1A1033 to void – page and panel backgrounds); Halo (radial, brand at 55% to transparent – ambient light behind glass, one per screen); Surface sheen (180deg, white 9% to 2% – the default card fill, light always from above); Stroke (a 1px gradient wrapper, white 35% to brand 35% to white 4% – edge lighting on featured cards).

04 ELEVATION – E0 flat, E1 resting card, E2 hover, E3 overlay, E4 brand glow. Print the exact shadow values. State the reason the scale exists: on a near-black base a shadow alone reads as almost nothing, so every level pairs a black shadow with a 1px hairline and a top inset highlight – the hairline is what separates surface from void. E4 is reserved for the primary action and focus. Never decorative.

05 RADIUS – 6, 10, 14, 20, 28, full. One tile each, with what it is for. Then the nesting rule, stated as a rule: inner radius equals outer radius minus padding. Concentric corners, never a tighter curve inside a looser one.

06 GLASS CARDS – three tiers, over one halo, so the difference is visible rather than asserted. Veil: 4% white fill, 12px blur, for panels laid directly on the halo. Standard: top-down sheen 10% to 4.5%, 20px blur at 140% saturation, visible inset highlight – the card you reach for. Featured: brand tint at 20%, 28px blur, 1px gradient stroke, outer glow; one per screen. Follow the three with a DO and a DON'T list, three each, written as instructions someone could break.

07 ASSEMBLED – one composed panel that uses every token at once and proves they survive contact: a workspace header, two stat tiles carrying real numbers, and one featured glass card as the upgrade prompt. If a token cannot earn its place here, delete it upstream.

08 FORM ELEMENTS – 44px targets, 10px radius. Text input at rest and focused, an input in its error state, a search field with an inline affix, a select, a textarea, checkboxes including a disabled one, radios, switches on and off, a slider, a segmented control, and six button variants. Affixes sit inside the field, never as separate boxes. State the focus treatment once and apply it to everything.

HOW THE PAGE ITSELF IS BUILT

Dark (#0A061A), single column, max-width 1180px, 96px of top padding, 112px between sections. One radial purple halo bleeding in from above the header and nowhere else. Header: a small brand-gradient tile beside a mono eyebrow reading VISUAL FOUNDATIONS · V1; a two-line display headline where the second line takes a white-to-lilac-to-magenta gradient text fill; one paragraph of intent, capped at 620px. Every section opens with a numbered rule row – mono number in brand purple, section name in Outfit, a right-aligned qualifier, a hairline beneath. Footer: a mono line naming the typefaces on the left and the four hex values on the right.

OUTPUT RULES

One HTML file. No build step, no framework, no external CSS. Fonts from Google Fonts. Static – no interactivity beyond what the browser gives you for free. Every specification printed in mono beside the thing it specifies. No lorem ipsum: every label, number and sentence is real content about this system. If a decision is genuinely open, print it as an open question rather than choosing quietly.

STAGE 05 · MOVE: STYLE · PART THREE

What came back

A single HTML page — every decision rendered as the thing it describes, with its specification printed in mono beside it.

VISUAL FOUNDATIONS · V1

Purple in the black. Light on the glass.

Gradients, elevation, radii and glass treatments built on a four-colour palette. Outfit carries every heading; Schibsted Grotesk carries every word you actually read.

04-DESIGN SYSTEM FOUNDATIONS.HTML · HEADER · THE HALO, THE TYPE, THE INTENT

01 Palette

4 tokens, no more



Void
#0A061A
Base. Reads as black, but there's purple in it.



Brand
#8A861A
The primary purple. This is the brand.



Magenta
#A855F7
#B440CB
Deeper magenta, used for accents.



Paper
#FFFFFF
Text.

text / primary
100% white

text / secondary
62% white

text / muted
42% white


stroke / hairline
white 10%

surface / raised
white 6%


01 PALETTE · FOUR COLOURS, AND THE FIVE TOKENS DERIVED FROM THEM

STAGE 05 · THE RESULT, CONTINUED

Decisions, rendered



02 TYPE · THE SPECIMEN LADDER, WITH THE RULES PANEL SET BESIDE IT




03 GRADIENTS · SIX, NAMED, EACH WITH ITS DECLARATION AND ITS PERMITTED USE

STAGE 05 • THE RESULT, CONTINUED

Depth and corners

04 Elevation


Shadow carries depth, glow carries state



EO

Flat


border 1px #FFF/8 • no shadow
Inline containers, list rows.



E1

Resting card


0 2px 6px -2px #000/50 • inset
0 1px 0 #FFF/9
The default card.



E2

Hover / raised


0 8px 24px -8px #000/65
Lift on hover, +1px translate.



E3

Overlay

0 24px 56px -16px #000/80
Menus, popovers, modals.



E4

Brand glow

0 12px 40px -10px #A855F7/75
Primary action and focus only.
Never decorative.

On a near-black base, shadow alone reads as almost nothing. Every level therefore pairs a black shadow with a 1px hairline and a top inset highlight — the hairline is what actually separates the surface from the void.

04 ELEVATION • E0 TO E4 • SHADOW CARRIES DEPTH, GLOW CARRIES STATE

05 Radius

6 • 10 • 14 • 20 • 28 • full

**6px**

Tags, checkboxes, inline code.

**10px**

Inputs, small buttons, icon tiles.

**14px**

Nested panels, list rows, media thumbs.

**20px**


Default card. The workhorse.

**28px**

Hero panels, modals, glass sheets.

**full**

Pills, avatars, toggles.




Nesting rule: inner radius = outer radius – padding. A 20px card with 10px padding holds a 10px child. Concentric corners, never a tighter curve inside a looser one.


05 RADIUS • SIX STEPS, AND THE CONCENTRIC NESTING RULE THAT GOVERNS THEM

STAGE 05 · THE RESULT, CONTINUED

Glass, assembled




06 GLASS · THREE TIERS OVER ONE HALO, WITH THE DO AND DON'T THAT GOVERNS THEM



07 ASSEMBLED · EVERY TOKEN AT ONCE · IF IT CANNOT EARN ITS PLACE HERE, DELETE IT UPSTREAM

STAGE 05 • THE DOCUMENT

One page, whole



04-DESIGN SYSTEM FOUNDATIONS.HTML • ALL EIGHT SECTIONS

DOCUMENT CREATED

04-Design System Foundations.html


[Open in Google Drive ↗](#)

WHY A PAGE AND NOT A DOCUMENT

A shadow described in prose is an opinion. A shadow rendered on the page is a decision. Building the foundations as HTML means every value is tested the moment it is written, and the build prompt inherits something that has already been seen to work rather than something that only reads well.

WHAT IT FEEDS

The build prompt, wholesale. No new colours, no new type sizes and no new radii are introduced downstream — if the page needs one, it comes back here first.



08 FORM ELEMENTS • 44PX TARGETS, ONE FOCUS TREATMENT

06

STAGE 06 · MOVE: SPECIFY

Build prompt

The handoff. One document a coding agent can execute end to end without asking a single question — because every question was already answered upstream.

WHAT IT PRODUCES

No document — the output is the site itself

WHAT IT FEEDS

The Review move — the impeccable audit, then iterate. All three builds are linked in 07, The page it produces.

05-Landing Page — Build Prompt.md

REPRODUCED IN FULL

Tessera Homepage — Build Prompt

Objective

Build the Tessera marketing homepage as a React application. One route, one scroll, ten sections, five conversion points. The audience is ML/platform engineers who deploy on it and the engineering leaders who approve it — both of whom leave when a page reads as marketing. Every claim on this page carries a number or a named mechanism. Nothing on it is decorative for its own sake.

Source documents: 02_Requirements.docx v1.0, 03_Content_Outline.docx v0.1, 04-Design System Foundations. Where this prompt and the content outline disagree, the outline wins on substance; this prompt wins on execution.

COMPANY

Tessera (tessera.run) — serverless GPU infrastructure for AI applications. Point it at an entry point or a Dockerfile; it provisions, scales and fails over across multiple clouds. Pay per second. No Kubernetes, no reservations, no rewrites.

Differentiator: 2-4 second cold starts via memory and GPU snapshotting, against a field sitting at 30+ seconds. 12+ GPU varieties. 99.999% uptime with multi-region failover. SOC 2, HIPAA, GDPR, ISO. gVisor isolation. Native OpenTelemetry.

TECH STACK

Vite 6 · React 19 · TypeScript, strict. Single route, static build, no SSR runtime.

Tailwind CSS v4 via @tailwindcss/vite. No tailwind.config.js — tokens are declared in CSS with @theme, which is the reason for picking v4 here: the design system becomes the token layer rather than a document someone has to remember to follow.

Motion (motion/react) for the five places that genuinely need it — the hero tile settle, the \$4 pinned pillar stepper (useScroll/useTransform over position:sticky), the cold-start comparison, the \$3 logo loop, and the \$6B testimonial columns. Everything else — hover, focus, reveal, form state — is a CSS transition. Do not reach for an animation library to do a hover state.

Fonts self-hosted via @fontsource-variable/outfit, @fontsource-variable/schibsted-grotesk, @fontsource-variable/jetbrains-mono. No external font requests, no layout shift on load.

No UI kit, no component library, no icon package. The four icons this page needs – chevron, copy, arrow, external – are hand-rolled inline SVG at 1.5px stroke. Installing lucide to draw an arrow is how a design system stops being a design system. No form library either; two email fields do not justify react-hook-form.

Prerender: a single marketing route on a CSR-only Vite SPA ships an empty shell to crawlers and to anyone the client forwards it to as a link preview. Add vite-react-ssg so the built output carries the real copy in the initial HTML. If the engagement later grows past this one page, the migration target is Next.js – say so in the build note rather than pretending Vite scales to a marketing site.

Project structure

```
src/  
  main.tsx  
  App.tsx  
  styles/tokens.css      @theme tokens, the six gradient utilities, font imports  
  content/homepage.ts    every string on the page, typed, one export per section  
  content/posts.ts       blog placeholders – PENDING OQ-04  
  content/proof.ts       placeholder marks and testimonials – PENDING OQ-07  
  content/limits.ts      per-slot character bounds from CPY-08, plus dev assertion  
  components/  
    Nav.tsx Hero.tsx HeroTiles.tsx ProofStrip.tsx LogoLoop.tsx  
    Pillars.tsx PillarStepper.tsx Workloads.tsx ColdStartScale.tsx  
    Security.tsx LeadForm.tsx Testimonials.tsx TestimonialColumn.tsx  
    Writing.tsx PostCard.tsx Closing.tsx Footer.tsx  
    ui/ Button · Field · GlassCard · Eyebrow · SectionHead · Pending · Mono  
  hooks/ useReveal.ts · useReducedMotion.ts · useLeadForm.ts · usePinnedSteps.ts  
  lib/events.ts          CTA destination + event-name registry, single source  
  lib/blocked.ts         every open question keyed by ref, rendered by <Pending/>
```

No string literals in components. All copy lives in content/homepage.ts so the copy deck (DEL-03) arrives as a file swap rather than a diff across fourteen components.

content/limits.ts holds the character bounds for every slot specified in this brief and asserts them at module load in dev – over-length throws with the slot name and the overage, under-length warns. CPY-08 exists so that real content cannot break the layout; enforce it in code rather than in review.

Token layer – styles/tokens.css

```
@import "tailwindcss";  
  
@theme {  
  --color-void: #0A061A;  --color-brand: #A855F7;  
  --color-magenta: #B440CB; --color-paper: #FFFFFF;  
  --color-raise-1: #150C2B; --color-raise-2: #1A1033;  
  --color-tint-1: #C77DFF; --color-tint-2: #D8B4FE; --color-tint-3: #E9D5FF;  
  --color-danger: #FB7185;  
  
  --font-display: "Outfit Variable", ui-sans-serif, system-ui;  
  --font-sans: "Schibsted Grotesk Variable", ui-sans-serif, system-ui;  
  --font-mono: "JetBrains Mono Variable", ui-monospace, monospace;  
  
  --radius-tag: 6px; --radius-control: 10px; --radius-panel: 14px;  
  --radius-card: 20px; --radius-sheet: 28px;  
  
  --shadow-e1: 0 2px 6px -2px rgb(0 0 0/.50), inset 0 1px 0 rgb(255 255 255/.09);  
  --shadow-e2: 0 8px 24px -8px rgb(0 0 0/.65);  
  --shadow-e3: 0 24px 56px -16px rgb(0 0 0/.80);  
  --shadow-e4: 0 12px 40px -10px rgb(168 85 247/.75);
```

```
--ease-brand: cubic-bezier(.32,.72,0,1);
}
```

The derived scale is opacity modifiers on those four colours and nothing else – text-paper/86, text-paper/62, text-paper/42, border-paper/10, bg-paper/6. If a value isn't reachable that way, it isn't in the system.

The six gradients ship as six @utility declarations – grad-brand, grad-brand-ext, grad-depth, grad-halo, grad-sheen, grad-stroke – and no component writes a raw linear-gradient(). Six utilities, six gradients, enforced by the absence of an alternative.

Arbitrary values (bg-[#8B5CF6], text-[19px], p-[37px]) are a build failure, not a style preference. If a value is needed and not in @theme, the token layer is what changes, and that change is a decision someone has to defend.

DESIGN SYSTEM – inherited, not negotiable

Palette – four tokens only

#0A061A Void Base. Reads as black, but there's purple in it.
#A855F7 Brand The primary purple.
#B440CB Magenta Deeper, used for accents.
#FFFFFF Paper Text.

Derived only: white at 100 / 86 / 68 / 62 / 42 / 35% for text tiers; white at 3-14% for surfaces and hairlines; #C77DFF / #D8B4FE / #E9D5FF as light brand tints for overlines and links; #150C2B / #1A1033 as raised background stops; #FB7185 and rgba(244,63,94,...) reserved for errors and destructive actions. No other hues. No green success ticks – a successful form state uses brand, not a new colour.

Typography

Outfit – every heading and every display numeral. 600 display, 500 default, 300 for oversized specimens only. Never 700. Never below 20px.
Schibsted Grotesk – all body copy, labels, buttons, form text. 17/1.62 body, 15/1.6 small.
Tabular figures in any table or stat row.
JetBrains Mono – eyebrows, code, key caps, certification labels, timing annotations. 11-12px, 0.16em tracking, uppercase for eyebrows.
Negative tracking scales with size: -3.5% display, -2.5% h1, -2% h2, 0 body.

Gradients – six, and only six

Brand linear-gradient(135deg,#A855F7,#B440CB)
Brand extended 115deg · #D8B4FE → #A855F7 42% → #B440CB 78% → #6B1F8A
Depth linear-gradient(180deg,#1A1033,#0A061A)
Halo radial-gradient(70% 90% at 50% 110%, #A855F7 55%, transparent)
Surface sheen linear-gradient(180deg,#FFF 9%,#FFF 2%) ← the default card fill
Stroke 1px padding wrapper · white 35% → brand 35% → white 4%

Elevation

E0 flat border 1px #FFF/8, no shadow – inline containers, list rows
E1 0 2px 6px -2px #000/50 · inset 0 1px 0 #FFF/9 – the default card
E2 0 8px 24px -8px #000/65 · +1px translate – hover
E3 0 24px 56px -16px #000/80 – menus, popovers, modals
E4 0 12px 40px -10px #A855F7/75 – primary action and focus only, never decorative
Every level pairs a black shadow with a 1px hairline and a top inset highlight. On near-black, shadow alone reads as nothing; the hairline is what separates surface from void.

Radius 6 tags/inline code · 10 inputs and small buttons · 14 nested panels and thumbs · 20 default card · 28 hero panels and glass sheets · full pills.
Nesting rule: inner radius = outer radius - padding. Concentric corners, never tighter inside looser.

Glass – three tiers over one halo

Veil #FFF/4 fill, light blur, border #FFF/10, r24
Standard sheen #FFF/10→4.5 · blur ~24px sat 140% · border #FFF/14 · inset 0 1px 0 #FFF/16
Featured tint #A855F7/20 · strong blur · 1px gradient stroke · glow #A855F7/70
One halo on the page. One brand glow per viewport. Never blur behind a scrolling list.
Never set body text below 62% white on glass. Cap glass at two layers deep.

Motion

All easing cubic-bezier(.32,.72,0,1). Scroll reveals are 8-12px translate plus opacity, no scale, no stagger longer than 60ms per item, and they fire once. Hover on cards is the E1→E2 transition and a 1px lift at 180ms. Everything obeys prefers-reduced-motion: the page must be complete, legible and operable with all motion removed.

Discrepancy to note: the foundations build prompt specified Instrument Sans; the delivered foundations sheet ships Schibsted Grotesk. Build with Schibsted Grotesk and flag it.

BRAND MARK – use the supplied geometry verbatim, do not redraw

The mark is five rounded squares rotated 45°, arranged as a mosaic – three across the top, two descending from the centre. Tesserae. Use the white lockup in the nav and the gradient mark in the footer and favicon. Paths, on a 139×125 viewBox for the mark and 520×125 for the lockup:

```
M11 7H29A4 4 0 0 1 33 11V29A4 4 0 0 1 29 33H11A4 4 0 0 1 7 29V11A4 4 0 0 1 11 7Z rotate(45 20 20)
M60 7H78A4 4 0 0 1 82 11V29A4 4 0 0 1 78 33H60A4 4 0 0 1 56 29V11A4 4 0 0 1 60 7Z rotate(45 69 20)
M109 7H127A4 4 0 0 1 131 11V29A4 4 0 0 1 127 33H109A4 4 0 0 1 105 29V11A4 4 0 0 1 109 7Z rotate(45 118 20)
M60 49H78A4 4 0 0 1 82 53V71A4 4 0 0 1 78 75H60A4 4 0 0 1 56 71V53A4 4 0 0 1 60 49Z rotate(45 69 62)
M60 91H78A4 4 0 0 1 82 95V113A4 4 0 0 1 78 117H60A4 4 0 0 1 56 113V95A4 4 0 0 1 60 91Z rotate(45 69 104)
```

Gradient fill: linearGradient x1=7 y1=7 x2=132 y2=118, stops #C77DFF → #A855F7 @0.55 → #B440CB.
Wordmark path is supplied in tessera-lockup-white.svg – inline it, do not re-typeset it.

THE HERO 3D MOMENT – concept is fixed, execution is open

Not a GPU chip. Not a server rack. Not an orbiting sphere of nodes – that is now a category cliché and it will be read as one.

The concept is state restoration: something reconstituting into form, which is literally what snapshotting does. The mark already is that idea – discrete tiles that resolve into a whole. Build the hero visual out of the tessera geometry: a field of rotated tiles held in scattered, unresolved positions that settle into the mark's five-tile arrangement over roughly 2-4 seconds, matching the claim in the H1. The settle happens once on load, and the duration is the argument. Tiles carry the brand-extended gradient and catch light from above.

Constraints: static fallback required (VIS-08). Reduced-motion alternative required – resolved state, no animation (VIS-09). All hero meaning survives with the visual absent (VIS-10). Hero text and CTAs must be legible and operable before the visual finishes loading (PRF-05). If the treatment forces a contrast compromise on the H1, the treatment changes, not the text (ACC-02). If it costs frames, it becomes static (VIS-07).

COPY – verbatim unless marked

Copy below is indicative-but-real, taken from the content outline. Use it as written. Do not substitute lorem ipsum anywhere, including in the blog module (CPY-01). Do not alter stated figures or invent new ones – the substance is client-supplied and authoritative.

Banned, in any casing, anywhere on the page (CPY-03):

scale effortlessly · built for AI · seamlessly · lightning-fast · effortless ·
enterprise-grade · request a demo · get started today · unlock your free trial ·

supercharge · unlock · revolutionise · game-changing

The H1 must carry a number (CPY-04). Named mechanisms – snapshotting, gVisor, OpenTelemetry, multi-region failover – travel with their claims and are never dropped for brevity (CPY-05).

REFERENCES

Three visual references and three component references. Visual references set energy and structure; they are not moodboards to match. Component references set interaction and layout; they are not code to install.

Reference A – dark observability long-scroll (attached)

The closest tonal sibling to the foundations: near-black base, single purple halo, glass product cards, long editorial scroll.

Take – the halo bleeding from behind the hero and never appearing again; the full-bleed band of hard numbers set at display scale with mono labels beneath; isometric line diagrams drawn in pure stroke on dark, used to explain mechanism rather than decorate; the section-opening eyebrow + rule; the density of a page that expects to be read.

Avoid – the orange accent; the repeating left-text / right-screenshot cadence, which is the single most tired pattern in infrastructure marketing; the light-on-white sections. Tessera is dark end to end. The one permitted inversion on the whole page is a single brand-gradient full-bleed moment at the closing CTA (§8), and nothing else.

Reference B – OneWhale (attached)

Take – the nerve of it. Oversized display type used as a structural device: a heading given an entire viewport with nothing else in it. Flat colour-field section breaks. Small mono metadata pinned to section corners. Space used as confidence rather than as padding.

Use this once, at §8. A page that shouts in four places is not confident, it is loud.

Avoid – the acid yellow; the consumer-crypto register; floating 3D app icons; type that overflows the viewport horizontally.

Reference C – griffin.com (attached, and see the live site)

The interaction model for §4, specified in full there.

Take – mono nav and mono eyebrows at wide tracking; the 01/02/03/04 numbered index; the thin vertical progress rail with ring markers marking position; hairline vertical rules dividing the viewport into a visible grid; exploded isometric layer diagrams in pure line.

Avoid – Griffin's serif display, the teal accent, and the beige pill CTA. Our display face is Outfit and our accent is brand purple.

Component references – port the pattern, do not install the package

§3 logo strip <https://reactbits.dev/animations/logo-loop>

§6B testimonials <https://21st.dev/@sshahaider/components/testimonials-columns-1>

§5 workloads <https://21st.dev/@meschacirung/components/features-9>

All three are copy-in sources, not libraries, and all three arrive styled as somebody else's product. Take the layout logic and the motion behaviour; rebuild the surface from our token layer. Specifically: features-9 ships depending on recharts, dotted-map and lucide-react – none of those enter this project. The one chart on the page is a single comparison bar and is hand-rolled inline SVG; the map, if used, is built from the tessera tile motif rather than a generic dot map; the icons are our four hand-rolled strokes. A reference that adds three dependencies to draw two shapes has not been ported, it has been pasted.

SECTIONS

1 · Navigation – capture point 1

Wordmark → home. Links: Product · Docs · Pricing · Blog · Company (4–9 chars each).

CTA A "Talk to us" – enterprise, higher emphasis, brand gradient fill, E4 glow.

CTA B "Start building" – self-serve, full button with hairline border. Not a text link, not subordinated (LED-04).

Sticky, becomes a glass Veil bar with hairline bottom after 40px of scroll. On mobile the collapsed nav keeps at least one CTA visible without opening the menu (RSP-06). Every CTA needs a defined href and a data-event attribute for tracking (LED-16). Both come from lib/events.ts – no CTA hardcodes its own destination, and every unknown destination is a typed placeholder against its open question rather than a "#".

2 · Hero – capture point 2, primary

H1 Your code on thousands of GPUs, cold-started in 2-4 seconds.

Sub Point Tessera at an entry point or a Dockerfile. We handle GPU provisioning, scaling and failover across multiple clouds.
No reservations. No rewrites. No cluster to run.

CTA Talk to us · Start building

Proof 12+ GPU varieties · 99.999% uptime · SOC 2, HIPAA, GDPR, ISO

H1 45-70 chars, max 2 lines desktop / 3 mobile. Subhead 140-220 chars, max 3 lines desktop / 5 mobile. Proof line 60-90 chars, wraps to 2 lines on mobile without orphaning a certification mark. Set the proof line in mono, muted, with hairline dividers.

Two alternate H1 directions live in content/homepage.ts as exported constants, switchable in one line for the direction review (DEL-01). Not visible states, not a toggle in the UI:

Alt 1 Two to four seconds from cold to serving.

Alt 2 Bring the container. We'll bring the fleet.

3 · Proof strip – BUILD, with placeholder marks

Status change from content outline v0.1, which omits this section. No logos, case studies or usage metrics have been supplied (OQ-07), so the module is built and populated with placeholder marks to hold the slot and prove the layout. Read the placeholder rules at the foot of this section before writing any of it – they are the condition on which this reversal stands.

Placed immediately below the hero fold. Continuous horizontal logo loop, per the reactbits LogoLoop reference: a single track duplicated end to end and translated on transform for a seamless wrap, paused on hover and on focus-within, with the leading and trailing edges masked to transparent so marks dissolve into the void rather than clipping at a hard edge. Duplicate count derives from measured track width, not a hardcoded x2 – the loop must not show a gap on an ultrawide viewport. Reduced motion freezes the track and renders a static, centred row.

Label Teams running production inference on Tessera:

Marks sit at a common optical height, not a common bounding box – set the height per mark. Render them at paper/42, lifting to paper/86 on hover. No colour, no card, no border, no drop shadow. A logo strip that needs chrome is a logo strip that does not have logos worth showing.

Placeholder rules – binding

Placeholder marks are invented wordmarks for companies that do not exist, set in Outfit with a simple geometric glyph. They must be plainly fictional: no real company names, no marks that resemble real marks, no near-misses on a real logo. Using an unlicensed real logo on a client's live site is a legal problem, not a content gap.

Every placeholder is registered in lib/blocked.ts against OQ-07 and rendered through <Pending/>, which in review builds carries a visible PLACEHOLDER marker above the strip. This section is the first item on the pre-launch checklist in the build note. It does not go public in this state. The outline's judgement still holds – with this audience thin proof is worse than none – and it is suspended for the direction stage only, so the client can see what the page becomes once OQ-07 is answered.

4 · The four pillars – comprehension only, no CTA inside

Eyebrow What Tessera does

H2 Four things that are usually your problem.

Standfirst Cold starts, capacity, code compatibility and visibility.

Tessera takes all four.

- 1 2-4 second cold starts / Low latency from the first request.
Containers launch in seconds using memory and GPU snapshotting for fast restores. Sudden bursts and scale-outs are handled automatically, without degrading performance. Figure: 2-4s
- 2 Elastic GPU scaling / No reservations, no lock-ins.
Instant access to thousands of GPUs across multiple clouds and regions. Workloads scale in real time – no capacity planning, no reservations, no infrastructure to manage.
- 3 Your code, your way / Bring your own code. We'll run it.
No rewrites, no decorators, no custom SDKs. Point Tessera at an entry point or a Dockerfile and the application runs exactly as it is – versioned, reproducible, ready to scale. Artefact: see below
- 4 End-to-end observability / Full visibility for every workload.
Complete visibility into every request: logs, metrics, scaling events and system performance, in real time. Native OpenTelemetry support plugs into the stack you already run.

Four identical icon-and-text cards are explicitly not acceptable (PIL-03). The pillars are not symmetrical in kind and the layout must exploit that: pillar 1 is a number and takes the largest typographic treatment on the page after the H1; pillar 2 is scale; pillar 3 is an artefact; pillar 4 is a data surface. Equal structural weight means equal hierarchy, not identical layout (PIL-01) – pillar 1 may carry more visual force through its figure without demoting the other three.

Interaction – the Griffin model

Take the scroll behaviour from griffin.com's "Banking infrastructure built to evolve" section, where 01 YOUR APP / 02 YOUR TECH / 03 OUR PLATFORM / 04 OUR BANK LICENSE advance as the reader scrolls. Four steps, four pillars, and the mapping is exact.

The section pins for four viewport-heights of scroll. Left column: a vertical progress rail with four ring markers and four mono indices – 01 COLD STARTS / 02 ELASTIC SCALING / 03 YOUR CODE / 04 OBSERVABILITY. The active step shows its heading at h2 scale, its subhead and its body; the three inactive steps collapse to index and label at paper/42. The rail fills between markers as the step advances. Right column: the pillar's own visual, which swaps as the step changes.

This interaction is the answer to PIL-03, not a decoration on top of it. Four identical cards are rejected because the pillars are not the same kind of thing; the swapping right panel is what lets each one have its own treatment inside one coherent frame. What appears there is different every time: pillar 1 is the 2-4s figure at the largest type on the page after the H1; pillar 2 is an isometric line diagram of multi-cloud, multi-region pooling, drawn in the Reference A / Griffin manner; pillar 3 is the code artefact; pillar 4 is a live-looking data surface of logs, metrics and scaling events.

It also serves PIL-01 better than a 2x2 grid does. Every pillar gets the identical footprint, the identical duration and the identical treatment in turn, where a grid always implies a reading order and therefore a ranking.

Transitions: the outgoing panel fades and drops 8px, the incoming rises 8px, 320ms on the brand ease, cross-faded not sequenced. The rail fill and the marker rings animate on transform. Built with position:sticky plus useScroll/useTransform from motion/react – no scroll-jacking, no hijacked wheel events, no smooth-scroll library. The reader's scroll speed is the reader's business; the section maps position to step and nothing else.

Degradations, all required

Reduced motion – no pin, no swap. All four steps render expanded and stacked, each with its visual, in document order.

Below md – no pin. Same stacked layout. Pinned scroll on a touch device with a dynamic browser chrome is a bug generator, and four viewport-heights of pin on a phone is a

hostage situation.

No JS / prerendered HTML – the stacked layout is what ships in the markup, and the pin is progressive enhancement applied on mount. This section must be fully readable in the prerendered output.

Keyboard – every step is reachable and readable without scrolling through the pin; the rail markers are buttons that jump to their step.

Pillar 3 artefact: the claim is about the absence of work, so prove it by showing how little there is (PIL-04). Render a real, runnable four-to-six-line Dockerfile or a two-line CLI invocation in a mono block with a working copy-to-clipboard control. The snippet lives in content/homepage.ts flagged PENDING ENGINEER CONFIRMATION via lib/locked.ts – this audience reads code as code, and a fabricated snippet will be spotted.

Constraints: H2 30-55 · standfirst 80-130 · pillar heading 18-32 single line at all breakpoints · subhead 30-48 · body 130-190, max 4 lines desktop.

5 · Workloads, and why cold start decides them – capture point 3, self-serve weighted

Eyebrow Where this matters

H2 The same two seconds decide three different workloads.

Standfirst Batch, realtime and voice put different demands on infrastructure. Cold start is what they have in common.

a Realtime voice ← lead position, strongest treatment

A voice application has a conversational budget of a few hundred milliseconds. Infrastructure that takes thirty seconds to wake has already lost the call. Restoring in 2-4 seconds is what makes voice viable on serverless at all.

b Large-scale batch

Thousands of jobs across thousands of GPUs, scheduled and scaled without a reservation. Spin up for the run, release on completion, pay for neither side of the gap.

c 12+ GPU chip varieties

Match the chip to the workload rather than the other way round. Twelve-plus varieties across multiple clouds, with no commitment to any one of them.

Qualifier These are three of many. If it runs in a container, it runs here.

[PENDING ENGINEERING CONFIRMATION – CAP-04]

This is a causal argument, not a feature list, and it is the page's one chance to demonstrate cause and effect rather than assert a benefit (CAP-03). It must be visually distinguishable from the four pillars by different layout logic, not just smaller type (CAP-02).

Layout – the features-9 model

Take the grid from the 21st.dev features-9 reference: a bento of unequal cells inside one bordered container, divided by hairlines rather than separated by gaps, each cell carrying its own visual rather than a shared icon treatment. That structure satisfies CAP-02 outright – §4 is a pinned stepper, §5 is a static bento, and no reader will confuse the two.

Map the three workloads onto unequal cells, because they are not equally important:

Realtime voice	the dominant cell, spanning the full width of the top row. Holds the cold-start comparison described below. This is the causal payload and it gets the most area on the page.
Large-scale batch	lower-left cell. Visual: the tessera tile motif tiled out to thousands, scaling and releasing – jobs, not decoration.
12+ GPU varieties	lower-right cell. Visual: a region field built from the tile motif, twelve marked. Not a dotted world map, and not the dotted-map package.

Hairline dividers at paper/10, container at r28, cells square to each other with no inner

radius except where the container's corners require it – concentric per the radius rule.
Cell padding 32px desktop, 20px mobile. Body copy in each cell holds to the character
bounds below; the visual takes whatever height remains.

The qualifier line sits below the container, full width, mono, paper/62 – outside the grid
so it reads as a caveat on the whole section rather than a fourth workload.

Below md the bento becomes a single column in the order voice → batch → GPUs. Voice keeps
its larger type and its comparison visual; it does not degrade to a peer of the other two.

Voice is the second motion moment: a time-based comparison – a 30-second cold start against
a 2–4 second one, drawn to scale on a shared axis – makes the argument better than prose
can. Annotate the conversational budget. Static and reduced-motion versions required; if it
costs performance it ships static (VIS-07).

Capture point 3, at section close

Prompt Test a cold start yourself.

CTA Start building Secondary: Read the docs

Constraints: H2 40–65 · standfirst 100–150 · workload heading 14–28 · body 150–230 ·
qualifier 50–80.

6 · Security and compliance – capture point 4, enterprise weighted

Placed before the closing block so it removes the objection rather than confirming a
decision already made (TRU-06). A footer badge strip is explicitly rejected (TRU-02,
STR-04): all four elements appear with their explanatory substance, never reduced to logos
or labels (TRU-03).

Eyebrow Running in regulated environments

H2 Isolation, residency and failover – with the mechanism attached.

Standfirst Compliance claims are easy to make. These are the specifics
behind ours.

a SOC 2 · HIPAA · GDPR · ISO

Built to meet strict security and privacy standards, providing a compliant
foundation for sensitive and regulated workloads.

[BLOCKED – OQ-06 / TRU-07 / CPY-06: certified vs compliant vs aligned vs in
progress carries legal weight and must be confirmed per certification.
Marks to be supplied by client. Render as typographic placeholders with the
correct footprint; do not draw approximations of certification marks.]

b Data residency

Deploy workloads into specific regions to meet regulatory or contractual
requirements. Data stays exactly where it needs to be, and you choose where
that is.

c Isolation

Every workload runs on gVisor in a hardened, isolated environment. Strong
container isolation without the performance cost that usually comes with it.

d 99.999% uptime

Multi-region failover. If a region or a cloud goes down, traffic routes to
the next best alternative inside your constraints. Figure: 99.999%

The uptime figure must read as a specific commitment with its mechanism attached (TRU-05) –
multi-region failover travels with the number, in the same visual unit, never as a footnote.
Certifications are represented visually as well as in text (TRU-04).

Capture point 4, immediately at section close – highest-intent moment on the page

Prompt Need this reviewed against your own requirements?

CTA Talk to us

Field Work email – the only mandatory field (LED-07); any additional field
must be justified against conversion cost (LED-08)

Consent [BLOCKED – OQ-10. Reserve the space with correct height and a

placeholder line; do not write consent copy.]

Micro We'll reply within one business day.

Design and render every form state (LED-11): default, focus, filled, loading, success, field-level error, submission failure. Focus is brand border plus a 3px rgba(168,85,247,0.25) ring. Success confirms capture and states what happens next and when (LED-13). Error copy is specific and actionable (LED-12) – "Something went wrong" is not acceptable; say what failed and what to do. Validation must accept long and uncommon TLDs and must not reject legitimate work addresses (LED-14). Mobile uses type="email", inputmode="email", autocomplete="email" (LED-19).

Constraints: H2 45–70 · standfirst 80–130 · element heading 12–28 · body 140–200.

6B · Customer proof – ADDED, not in content outline v0.1

Placed between §6 and §7, so that named technical proof lands immediately after the trust argument and immediately before the credibility of the writing. It carries no CTA: capture point 4 is thirty pixels above it and a second ask here competes with the one that matters.

Eyebrow From teams running it

H2 What the engineers who deploy on it say.

Standfirst Not the procurement quote. The one from the person on call.

Three vertically scrolling columns, per the 21st.dev testimonials-columns-1 reference: column 1 always visible, column 2 from md, column 3 from lg, each looping at a different duration – roughly 15s / 19s / 17s – so the columns never lock into a visible rhythm. Vertical mask on the container, linear-gradient(to bottom, transparent, black 25%, black 75%, transparent), so cards enter and leave through the void rather than at a cut line. Pause on hover and on focus-within.

Card: quote in Schibsted Grotesk 17/1.62 at paper/86; below it, a rule; then name at paper/100, role and company at paper/62 in 15px. Glass Standard fill, r20, E1. No portrait photographs – invented faces are worse than no faces, and a monogram tile in the brand gradient at 40px does the same job honestly. No star ratings, no quotation-mark glyphs set at 80px, no company logos inside the cards.

Register: these are engineers, so the quotes read like engineers. Specific, unenthusiastic, carrying a number or a mechanism – a migration that took an afternoon, a p95 that moved, a Kubernetes cluster nobody has to run any more. Nothing that reads as a marketing quote. Nine to twelve cards, 140–260 characters each.

Reduced motion – columns freeze and the section renders as a static three-column grid of the first six cards, mask removed. Below md – one column, static, four cards, no loop. Never animate the columns while a card inside them has focus.

Placeholder rules – binding, same as §3

Every quote, name, role and company here is invented, and the §3 placeholder rules apply in full: plainly fictional companies, no real names, no real people, no resemblance to either. Registered in lib/blocked.ts against 0Q-07, rendered through <Pending/> with a visible PLACEHOLDER marker in review builds, and second on the pre-launch checklist. A fabricated testimonial on a live page is a different category of problem from a placeholder logo – it is a claim about a person who did not make it. This section ships to the client as a layout demonstration and to nobody else.

Constraints: H2 40–65 · standfirst 60–100 · quote 140–260 · name 4–28 · role and company 12–48.

7 · Engineering writing

This must not read as an afterthought appended to the bottom of the page; it needs deliberate design weight (BLG-03, STR-05).

Eyebrow From the team

H2 Benchmarks, architecture and things that broke.

Standfirst We publish what we measure, including the parts that didn't work.

Link All writing →

Three to four most recent posts (BLG-02). Each entry: title (40-95 chars, 2-line allowance, truncate at 2 lines with ellipsis per BLG-07), category (8-20 chars – Benchmark · Architecture · Customer build), standardised date, link, optional thumbnail that degrades gracefully when absent (BLG-05).

Build every state (BLG-08, STR-06): 3-4 posts is the designed default; 2 posts adapts to a two-column or centred pair with no visible gap; 1 post becomes a single featured treatment; 0 posts hides the section entirely without breaking the sections either side; no thumbnail becomes a designed typographic card, not a defaulted one. All five are reachable in the running app via a ?blog= query param (4 · 3 · 2 · 1 · 0 · nothumb) read in Writing.tsx – a review affordance, not a UI control, and stripped from the production build.

Write four real placeholder posts in the correct register with realistic title lengths – one at the 95-character ceiling to prove truncation. Mark them PENDING OQ-04. Route to the blog index even though the index is out of scope (BLG-06).

8 · Closing conversion – capture point 5

H2 Two ways to start.

Enterprise – primary

Talk to us

Bring your workload, your compliance requirements and your worst-case traffic. We'll tell you what it costs and how it behaves.

Field: Work email · Consent: [OQ-10] · CTA: Talk to us

Micro: We'll reply within one business day.

Self-serve – secondary but not subordinated

Start building

Deploy a container and watch it cold-start. No card, no call.

CTA: Start building

Micro: [BLOCKED – OQ-05: free tier / credits / waitlist wording]

Both paths visible without forcing a wrong choice (LED-04). Enterprise carries greater visual weight per commercial priority; the self-serve path stays a full, unhedged route – not a text link under the form. Constraints: H2 15-35 · path heading 10-20 · body 90-170 · microcopy 30-60.

9 · Footer

Wordmark plus descriptor – Serverless infrastructure for AI applications.

Columns: Product · Pricing · Docs · Changelog / Company · Blog · Careers · Contact / Security · Compliance · Status · Trust centre.

Certification marks repeat here; their presence in the footer never substitutes for §6 (TRU-01, TRU-02). Legal row: copyright, privacy policy, terms, cookie preferences (TEC-08). Status indicator if a status page exists.

CONVERSION LOGIC

Five capture points: nav (both paths), hero (both paths, primary), §5 close (self-serve weighted), §6 close (enterprise weighted), §8 (both paths). No capture point inside §4 – it is a comprehension block and interrupting it costs more than it captures – and none in §3 or §6B, which are evidence, not asks. §6B in particular sits directly beneath capture point 4; a second CTA there competes with the highest-intent moment on the page.

Every CTA carries an href and a data-event attribute, both resolved from lib/events.ts. That file is the deliverable as much as the page is: one typed record per CTA carrying its section, its intended destination, its event name, and its blocking open question where the destination is unknown (LED-16, OQ-05, OQ-12). Reproduce it as a table in the build note.

ACCESSIBILITY AND PERFORMANCE

WCAG 2.1 AA. Body text over any rendered surface clears 4.5:1 measured against the actual composited backdrop, not the flat base (ACC-02). Visible focus on every interactive element – brand border plus 3px ring, never outline:none without replacement. Semantic landmarks, one h1, logical heading order, real button and form elements. Full keyboard path through every CTA and form. prefers-reduced-motion removes all motion including the hero settle and the voice comparison, leaving both in resolved states.

Hero text and CTAs interactive before any visual asset resolves (PRF-05) – HeroTiles is lazy and the hero is complete without it. No layout shift on font load: preload the display and body subsets, size-adjust the fallback stack. No blur behind scrolling content.

Budget: under 150KB of JS gzipped for the whole route, LCP under 2.0s on a throttled 4G/mid-tier mobile profile, CLS under 0.02. Report the actual numbers in the build note – if a motion moment breaks the budget it ships static, per VIS-07, and that is not a negotiation.

RESPONSIVE

Mobile-first, using Tailwind's default sm/md/lg/xl scale – do not redefine the breakpoints to hit exact device widths. Verify at 1440 / 1024 / 768 / 390.

Max width 1180px, 48px gutters desktop, 20px mobile, 112px between sections desktop and 72px mobile. Sections reflow rather than shrink:

§3 loop continues, marks scale down, mask width reduces with the gutter
§4 pin released below md – four steps stacked and expanded, in document order
§5 bento collapses to one column, voice → batch → GPUs, voice keeps its scale
§6 security grid 2×2 then single
§6B three columns → one, static, four cards, no loop
§7 3-4 posts → 2 → 1 per the state rules, never a visible gap

Test the H1, the hero proof line, the longest testimonial and the longest blog title at 390px – none may orphan a word or a certification mark. Test §4 at 768px specifically, which is where a pin is most likely to be wrongly retained.

CONSTRAINTS

Tailwind utilities only, from the @theme token layer. No arbitrary values, no inline style props except where a value is genuinely computed at runtime – the hero tile transforms and the cold-start bar widths, and nowhere else. No CSS-in-JS, no styled-components, no second styling system arriving alongside Tailwind.

Flex/grid with gap for every sibling group; no margin stacking to fake spacing. Semantic elements before divs. One component per section, no component over ~150 lines – if a section needs more, it has a sub-component it hasn't been given yet.

No emoji. No stock imagery. No portrait photography, invented or otherwise. No SVG illustration beyond the mark geometry, the hero tile field, and the mechanism diagrams in §4 and §5 – which are line drawings built from the tile motif, not illustration. No invented colours and no invented figures, ever.

Dependencies are the three fontsource packages, motion, and vite-react-ssg. Nothing else without a written reason in the build note. The component references are ported by hand: no recharts, no dotted-map, no lucide-react, no marquee package, no shadcn registry pull.

Invented content – the one carve-out, and its boundary

Placeholder proof in §3 and §6B is permitted for the direction stage, against the outline's own judgement, so the client can see the assembled page. It is bounded: plainly fictional companies and people, no real names, no resemblance to real marks or real individuals, no invented usage metrics anywhere on the page. Every instance registered in lib/blocked.ts against OQ-07, rendered through <Pending/>, carrying a visible PLACEHOLDER marker in review builds. Figures stay untouched – 2-4s, 12+, 99.999% and the certification list are client-supplied and are not placeholder material.

Every blocked item renders at its correct footprint through <Pending/>, reading its ref and label from lib/blocked.ts, so nothing re-layouts when the answers arrive: OQ-04 blog source, OQ-05 self-serve destination and microcopy, OQ-06 certification wording, OQ-07 proof strip and testimonials, OQ-10 consent copy, OQ-12 analytics taxonomy, plus the pillar 3 artefact and the CAP-04 qualifier awaiting engineering. In production the markers render as neutral placeholders; in dev they carry their reference visibly.

Ships clean: tsc --noEmit passes, vite build passes with no warnings, no console output at runtime, no unused dependency in package.json.

DELIVER

The running app, then a short build note covering:

- which H1 direction is live and why
- the fallback behaviour of every motion moment under prefers-reduced-motion, below md, and in the prerendered no-JS output
- the CTA destination and event table from lib/events.ts
- measured JS weight, LCP and CLS against the budget
- the Instrument Sans / Schibsted Grotesk discrepancy flagged against the foundations
- two structural changes made against content outline v0.1, for sign-off: §3 built with placeholders rather than omitted, and §6B added between §6 and §7
- and a pre-launch checklist, first two items fixed: replace §3 placeholder marks with real logos or delete the section; replace §6B placeholder testimonials with real, attributed, permissioned quotes or delete the section. Neither ships public as built. Then: OQ-05 destinations, OQ-06 certification wording, OQ-10 consent copy, OQ-12 events, the pillar 3 artefact, and the CAP-04 qualifier.

07 • THE OUTPUT

The page it produces

Three builds, in the order they happened. The first is what the pipeline produced unassisted; the second is the same page after the impeccable audit; the third after the header was tuned by hand.

01 First output

Straight out of the build prompt — roughly seventy per cent right, which is the correct expectation.

tessera-homepage-three.vercel.app ↗

02 After the impeccable audit

The Review move applied: critique against the agreed documents, then iterate.

tesseraimpeccable.vercel.app ↗

03 After fine-tuning the header

The last pass, done by hand — the part taste is actually for.

tesserafinal.vercel.app ↗

07 • THE ARGUMENT, DRAWN

Why the number wins

Voice is the causal payload: a conversational budget of a few hundred milliseconds against a field that takes thirty seconds to wake. Drawn on a shared axis, the comparison makes the argument better than prose can.




FIG 3 • COLD START, TO SCALE

How the live page gets judged

LENS

WHAT IT MEANS

The outcome

Lead generation — self-serve signups and qualified conversations, without obstructing the traffic that goes straight to the docs and never talks to anyone.

Five seconds

Clear positioning. A reader should be able to say what this company does, and who it is not for, within five seconds. Every lead lost before that point is lost silently.

Personalised

What is in it for me, specifically. The voice buyer, the LLM buyer and the person running the security review want three different answers on the same page.

08 · FIELD NOTES

Going live, and what to expect

HOST IT PROPERLY

Vercel · Netlify · Cloudflare Pages

Push to a repo, connect once, every commit deploys with a preview URL for the client. Free tiers cover a marketing site comfortably. SSL, CDN and rollbacks are included and unconfigured.

WHERE THEY ALREADY PAY

GoDaddy · Hostinger · cPanel

A static build uploads and works. You give up preview deploys, instant rollback and edge caching, and you take on manual uploads. Fine for a brochure site — a real cost on anything you will iterate.

THE EXPECTATION TO SET ON THE FIRST BUILD

Six documents go in. Expect the first cut to be roughly **seventy per cent right** — that is the correct expectation, not a disappointment. The remaining thirty per cent is the Review move: the impeccable audit, the critique, the iteration.

The whole point of the pipeline is that this thirty per cent gets argued against documents everyone has already agreed to, rather than against taste alone. That is the difference between a revision round and a re-do.

THE ONE THING TO CARRY BACK

None of this is prompting technique. It is the studio process you already run, written down at a level of precision that a model can execute — which turns out to be the same level of precision a new hire would have needed anyway.



THANK YOU

Run the pipeline on your next brief

Tell us what was useful and what was not — it changes the next session more than you would think. Every question asked in the chat gets an answer, including the ones we could not get to live, within three to five working days, together with the recording.